

# PROYECTO VISIÓN POR COMPUTACIÓN CIFAR 10

Nicolás Goldschwartz

# VERSIÓN 1

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential\_1"

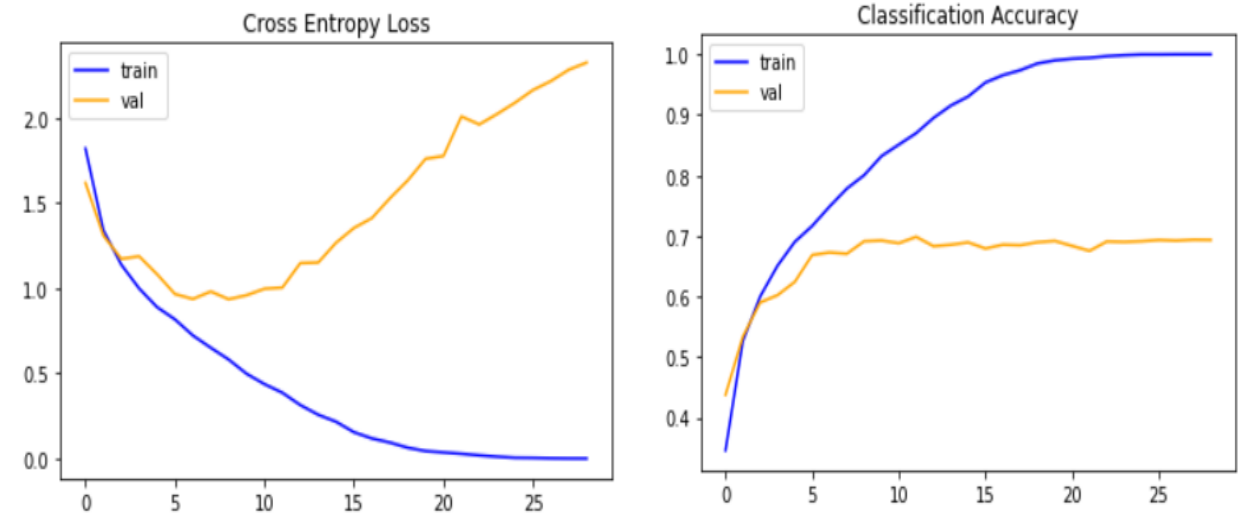
Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_1 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	36992
conv2d_5 (Conv2D)	(None, 16, 16, 256)	295168
flatten_1 (Flatten)	(None, 65536)	0
dense_2 (Dense)	(None, 32)	2097184
dense_3 (Dense)	(None, 10)	330

```
=====
Total params: 2,430,570
Trainable params: 2,430,570
Non-trainable params: 0
```

```
[34] history = model.fit(x_train_scaled, y_train, epochs=30,
                        callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                        batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
[36] _, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
      print('> %.3f' % (acc * 100.0))
```

> 68.860



Partimos desde el proyecto 0 con un modelo con una accuracy de 56.730. Para incrementar la precisión vamos a empezar por añadir a la arquitectura dos capas de convoluciones, una de 128 neuronas y otra de 256 para aumentar el barrido de las imágenes. El resultado ha sido muy satisfactorio , la accuracy ha aumentado más de 10 puntos hasta situarse en 68.860.

## VERSIÓN 2

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential\_2"


Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_2 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_7 (Conv2D)	(None, 16, 16, 128)	36992
conv2d_8 (Conv2D)	(None, 16, 16, 256)	295168
conv2d_9 (Conv2D)	(None, 16, 16, 128)	295040
conv2d_10 (Conv2D)	(None, 16, 16, 256)	295168
flatten_2 (Flatten)	(None, 65536)	0
dense_4 (Dense)	(None, 32)	2097184
dense_5 (Dense)	(None, 10)	330

```
=====
Total params: 3,020,778
Trainable params: 3,020,778
Non-trainable params: 0
```

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 68.430



Para esta segunda versión y siguiendo la línea de los resultados anteriores decidimos añadir otro paquete de dos convoluciones iguales que el anterior de 128 y 256 neuronas respectivamente. La accuracy ha disminuido un poco en vez de aumentar, resultando en un valor de 68.430. Es posible que el no incremento de accuracy se deba al hecho de no haber introducido ninguna capa maxpooling después de las capas de convolución. Las gráficas de función de error y de precisión no muestran bondad del modelo, sus líneas de entrenamiento y validación están muy distanciadas entre ellas en ambas gráficas.

# VERSIÓN 3

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential\_3"


Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_3 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_12 (Conv2D)	(None, 16, 16, 128)	36992
conv2d_13 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_4 (MaxPooling 2D)	(None, 8, 8, 256)	0
conv2d_14 (Conv2D)	(None, 8, 8, 128)	295040
conv2d_15 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_5 (MaxPooling 2D)	(None, 4, 4, 256)	0
flatten_3 (Flatten)	(None, 4096)	0
dense_6 (Dense)	(None, 32)	131104
dense_7 (Dense)	(None, 10)	330

```
=====
Total params: 1,054,698
Trainable params: 1,054,698
Non-trainable params: 0
```

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 73.960



Para corregir el poco incremento de accuracy de la anterior arquitectura vamos a introducir dos capas maxpooling, una después de cada paquete de convoluciones con la idea de reducir la dimensión espacial, es decir el ancho por alto de la imagen del volumen de entrada para la siguiente capa. El resultado ha dado lugar a un incremento de la accuracy de unos 5 puntos llegando a 73.960. Las gráficas de error y precisión del modelo siguen igual que en la versión anterior.

# VERSIÓN 4

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(256, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

Model: "sequential\_4"


Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_6 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_17 (Conv2D)	(None, 16, 16, 128)	36992
conv2d_18 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_7 (MaxPooling 2D)	(None, 8, 8, 256)	0
conv2d_19 (Conv2D)	(None, 8, 8, 128)	295040
conv2d_20 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_8 (MaxPooling 2D)	(None, 4, 4, 256)	0
flatten_4 (Flatten)	(None, 4096)	0
dense_8 (Dense)	(None, 32)	131104
dense_9 (Dense)	(None, 128)	4224
dense_10 (Dense)	(None, 256)	33024
dense_11 (Dense)	(None, 10)	2570

=====  
Total params: 1,094,186  
Trainable params: 1,094,186  
Non-trainable params: 0

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                   callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                   batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 73.540



Se introducen dos capas dense tras el flatten de 128 y 256 neuronas de entrada y salida respectivamente para aumentar las conexiones en la zona fully conected y obtener así una mejora en el área de clasificación. Finalmente la accuracy ha resultado un poco inferior a la versión anterior situándose en 73.540. Al parecer este incremento de capas tras el flatten no han sido de ayuda. Tampoco hay mejoras en las gráficas de error y precisión.

## VERSIÓN 5

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(32, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

Model: "sequential"

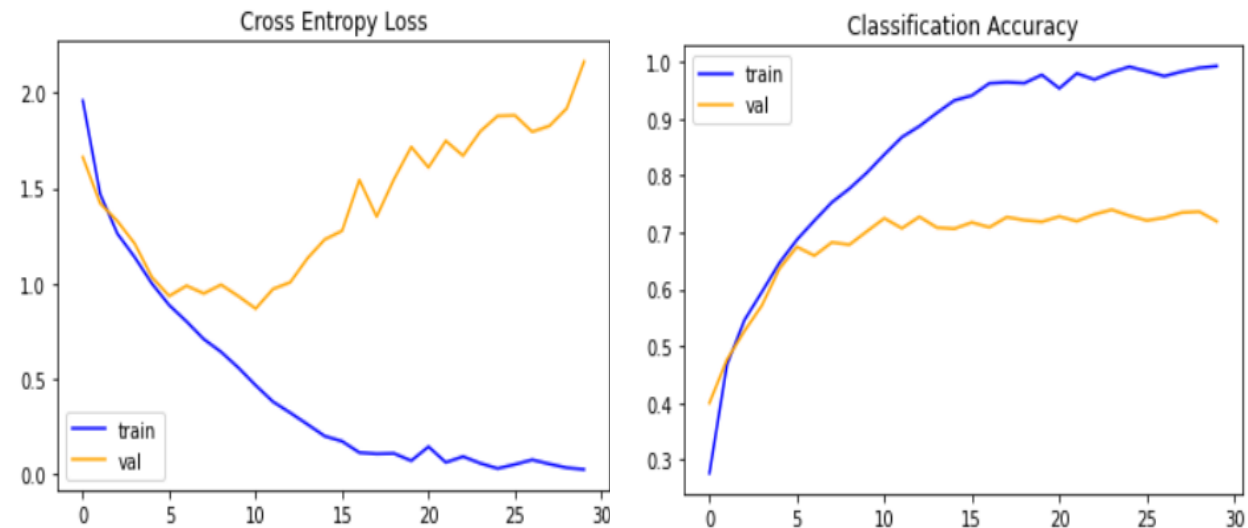
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 128)	36992
conv2d_2 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_3 (Conv2D)	(None, 8, 8, 128)	295040
conv2d_4 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 256)	1048832
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 32)	4128
dense_3 (Dense)	(None, 10)	330

```
=====
Total params: 2,009,450
Trainable params: 2,009,450
Non-trainable params: 0
```

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 71.580



Se realizan cambios en el orden de neuronas entre la capa flatten y la de softmax de modo que se pasa de orden creciente a decreciente, empezando con una capa de 256 otra de 128 y una final de 32 neuronas . Esta modificación ha arrojado una accuracy inferior a la de la versión anterior perdiendo dos puntos y quedando en 71.580. El overfitting y la validación de error siguen sin mejoras.



# VERSIÓN 6

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(64, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

Model: "sequential\_1"

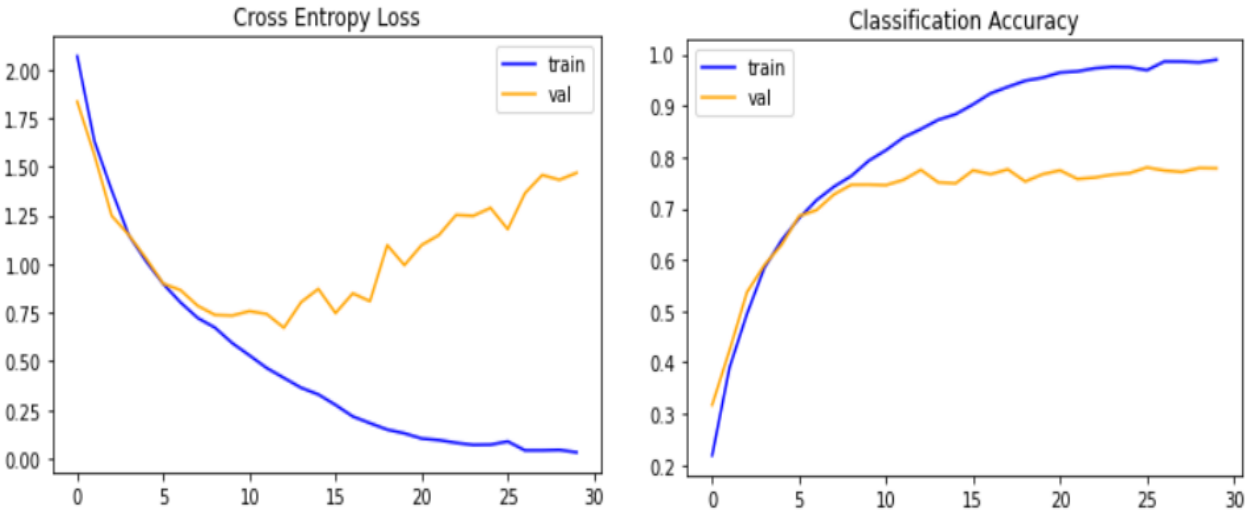
Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 32, 32, 32)	896
conv2d_6 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_3 (MaxPooling 2D)	(None, 16, 16, 64)	0
conv2d_7 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_8 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_4 (MaxPooling 2D)	(None, 8, 8, 256)	0
conv2d_9 (Conv2D)	(None, 8, 8, 128)	295040
conv2d_10 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_5 (MaxPooling 2D)	(None, 4, 4, 256)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_4 (Dense)	(None, 32)	131104
dense_5 (Dense)	(None, 64)	2112
dense_6 (Dense)	(None, 128)	8320
dense_7 (Dense)	(None, 10)	1290

=====  
Total params: 1,121,450  
Trainable params: 1,121,450  
Non-trainable params: 0

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                   callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                   batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 77.310



Se introducen varias correcciones , en primer lugar se añade una capa de convolución de 64 neuronas al primer paquete, quedando finalmente tres capas de convolución seguidos de un maxpooling cada una en segundo lugar se modifica el número de neuronas y su orden entre el flatten y la última capa se establece un orden ascendente de 32, 64 y 128 neuronas, la idea es que mejore en la precisión en la clasificación de las imágenes . El resultado ahora sí es notable , la accuracy ha mejorado sustancialmente hasta alcanzar los 77.310.

# VERSIÓN 7

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(64, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(512, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 32, 32, 32)	896
conv2d_12 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_6 (MaxPooling 2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
conv2d_13 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_14 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_7 (MaxPooling 2D)	(None, 8, 8, 256)	0
dropout_1 (Dropout)	(None, 8, 8, 256)	0
conv2d_15 (Conv2D)	(None, 8, 8, 128)	295040
conv2d_16 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_8 (MaxPooling 2D)	(None, 4, 4, 256)	0
dropout_2 (Dropout)	(None, 4, 4, 256)	0
flatten_2 (Flatten)	(None, 4096)	0
dense_8 (Dense)	(None, 32)	131104
dense_9 (Dense)	(None, 64)	2112
dense_10 (Dense)	(None, 128)	8320
dense_11 (Dense)	(None, 512)	66048
dense_12 (Dense)	(None, 10)	5130


=====

Total params: 1,191,338  
Trainable params: 1,191,338  
Non-trainable params: 0

```
history = model.fit(x_train_scaled, y_train, epochs=30,
callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_ , acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 80.100



En esta versión se introducen dos cambios principales uno para corregir el overfitting y otro para mejorar la precisión en la clasificación, para el primero se introducen tres capas de dropout con un rate de 0.4 tras el maxpooling de cada paquete de convoluciones. Para la segunda modificación se añade una capa dense de 514 neuronas en último paquete tras el flatten. Los cambios introducidos han aumentado la accuracy llegando a un 80.100 y los dropouts han mejorado notablemente el overfitting.



# VERSIÓN 8

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(64, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(512, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential\_3"

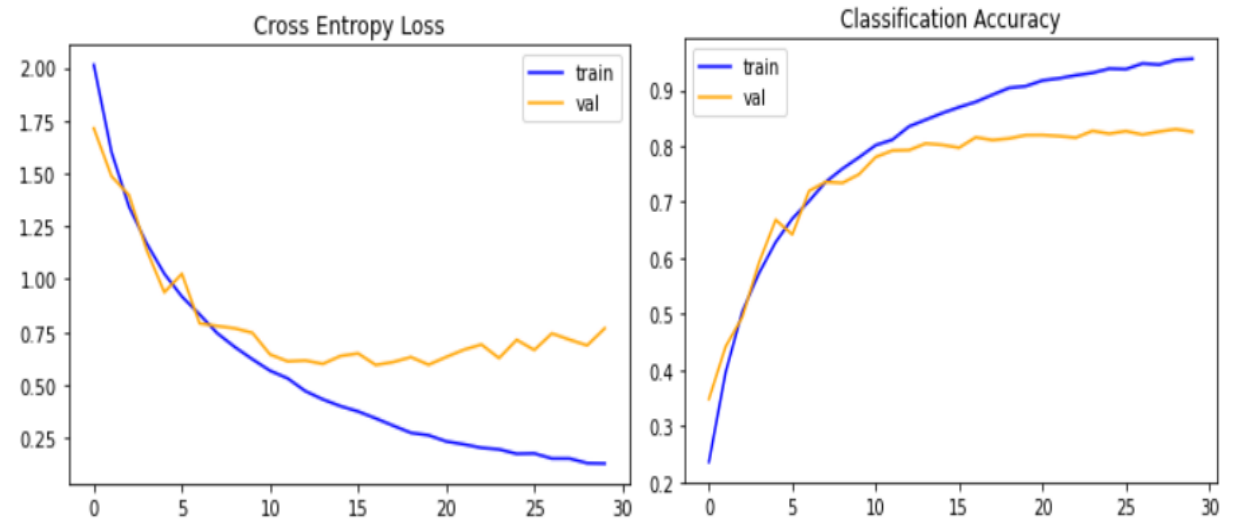
Layer (type)	Output Shape	Param #
conv2d_17 (Conv2D)	(None, 32, 32, 32)	896
conv2d_18 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_9 (MaxPooling 2D)	(None, 16, 16, 64)	0
dropout_3 (Dropout)	(None, 16, 16, 64)	0
conv2d_19 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_20 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_10 (MaxPoolin g2D)	(None, 8, 8, 256)	0
dropout_4 (Dropout)	(None, 8, 8, 256)	0
conv2d_21 (Conv2D)	(None, 8, 8, 512)	1180160
conv2d_22 (Conv2D)	(None, 8, 8, 1024)	4719616
max_pooling2d_11 (MaxPoolin g2D)	(None, 4, 4, 1024)	0
dropout_5 (Dropout)	(None, 4, 4, 1024)	0
flatten_3 (Flatten)	(None, 16384)	0
dense_13 (Dense)	(None, 32)	524320
dense_14 (Dense)	(None, 64)	2112
dense_15 (Dense)	(None, 128)	8320
dense_16 (Dense)	(None, 512)	66048
dense_17 (Dense)	(None, 10)	5130

Total params: 6,894,122  
Trainable params: 6,894,122  
Non-trainable params: 0

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 81.940



Se introducen más neuronas al último paquete de convoluciones incrementando en 512 y 1024 respectivamente, de manera que la arquitectura del modelo va incrementando el número de neuronas desde la primera a la última convolución, empezando en 32 y finalizando en 1024. La accuracy ha mejorado casi dos puntos situándose en 81.940 pero el overfitting ha aumentado en relación a la versión anterior

# VERSIÓN 9

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.BatchNormalization())
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dense(512, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential\_12"


Layer (type)	Output Shape	Param #
conv2d_72 (Conv2D)	(None, 32, 32, 32)	896
conv2d_73 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_36 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_36 (Dropout)	(None, 16, 16, 64)	0
conv2d_74 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_75 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_37 (MaxPooling2D)	(None, 8, 8, 256)	0
dropout_37 (Dropout)	(None, 8, 8, 256)	0
conv2d_76 (Conv2D)	(None, 8, 8, 512)	1180160
conv2d_77 (Conv2D)	(None, 8, 8, 1024)	4719616
batch_normalization_16 (Batch Normalization)	(None, 8, 8, 1024)	4096
max_pooling2d_38 (MaxPooling2D)	(None, 4, 4, 1024)	0
dropout_38 (Dropout)	(None, 4, 4, 1024)	0
flatten_12 (Flatten)	(None, 16384)	0
dense_60 (Dense)	(None, 32)	524320
dense_61 (Dense)	(None, 128)	4224
dense_62 (Dense)	(None, 256)	33024
dense_63 (Dense)	(None, 512)	131584
dense_64 (Dense)	(None, 10)	5130

-----  
Total params: 6,990,570  
Trainable params: 6,988,522  
Non-trainable params: 2,048

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                   callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                   batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 82.320



Modifico el número de neuronas de dos capa tras el flatten pasando a 32, 128, 256 y 512. Además añado una capa de batch normalization tras el tercer grupo de convoluciones con la idea de que normalice los mini batches para evitar que tengan muchas diferencias entre ellos con el propósito de que sean mas similares los datos que llegan a los filtros de modos que la red entrene mejor y el modelo adquiriera mayor robustez. El resultado ha sido muy satisfactorio aumentando la acuraccy hasta 82.320

# VERSIÓN 10

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.BatchNormalization())
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dense(512, activation='relu'))
model.add(ks.layers.Dense(1024, activation='relu'))
```

Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_3 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 256)	0
dropout_1 (Dropout)	(None, 8, 8, 256)	0
conv2d_4 (Conv2D)	(None, 8, 8, 512)	1180160
conv2d_5 (Conv2D)	(None, 8, 8, 1024)	4719616
batch_normalization (Batch Normalization)	(None, 8, 8, 1024)	4096
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 1024)	0
dropout_2 (Dropout)	(None, 4, 4, 1024)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 32)	524320
dense_1 (Dense)	(None, 128)	4224
dense_2 (Dense)	(None, 256)	33024
dense_3 (Dense)	(None, 512)	131584
dense_4 (Dense)	(None, 1024)	525312
dense_5 (Dense)	(None, 10)	10250

=====  
Total params: 7,521,002  
Trainable params: 7,518,954  
Non-trainable params: 2,048

```
history = model.fit(x_train_scaled, y_train, epochs=30,
                   callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                   batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 83.310



Al Introducir una capa más de dense de 1024 neuronas tras el flatten y situada en última posición, la acuraccy ha mejorado casi un punto hasta situarse en 83.310 a su vez las gráficas muestran unos resultados muy similares a las versiones anteriores

# VERSIÓN 11

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.BatchNormalization())
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dense(512, activation='relu'))
model.add(ks.layers.Dense(1024, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```

Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_3 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 256)	0
dropout_1 (Dropout)	(None, 8, 8, 256)	0
conv2d_4 (Conv2D)	(None, 8, 8, 512)	1180160
conv2d_5 (Conv2D)	(None, 8, 8, 1024)	4719616
batch_normalization (Batch Normalization)	(None, 8, 8, 1024)	4096
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 1024)	0
dropout_2 (Dropout)	(None, 4, 4, 1024)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 32)	524320
dense_1 (Dense)	(None, 128)	4224
dense_2 (Dense)	(None, 256)	33024
dense_3 (Dense)	(None, 512)	131584
dense_4 (Dense)	(None, 1024)	525312
dense_5 (Dense)	(None, 10)	10250

-----  
Total params: 7,521,002  
Trainable params: 7,518,954  
Non-trainable params: 2,048

```
history = model.fit(x_train_scaled, y_train, epochs=100,
callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 84.780



Mantengo la misma arquitectura que en la versión anterior y aumento las epochs a 100, la idea es que la red aprenda mas tiempo. El resultado ha sido satisfactorio y la accuracy ha aumentado hasta 84.780 .

Las gráficas parecen indicar que tanto validación error como validación de precisión del modelo se mantienen con cierto overfitting.



# VERSIÓN 12

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(64, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.BatchNormalization())
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(512, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(1024, activation='relu'))

model.add(ks.layers.Dense(10, activation='softmax'))
```


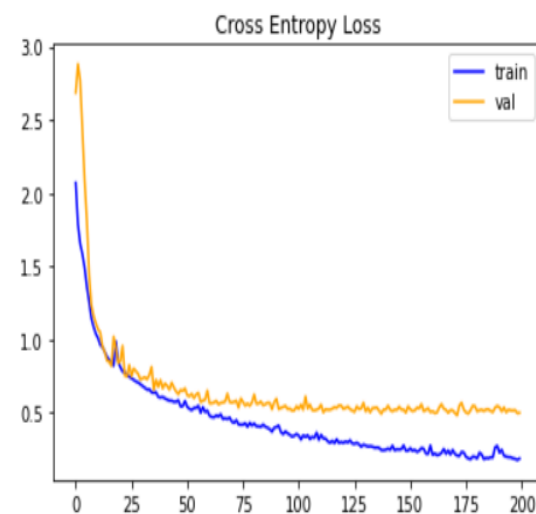
```
model.summary()

=====
conv2d_43 (Conv2D)          (None, 32, 32, 32)      896
conv2d_44 (Conv2D)          (None, 32, 32, 64)     18496
max_pooling2d_21 (MaxPoolin (None, 16, 16, 64)      0
g2D)
dropout_26 (Dropout)        (None, 16, 16, 64)      0
conv2d_45 (Conv2D)          (None, 16, 16, 128)    73856
conv2d_46 (Conv2D)          (None, 16, 16, 256)   295168
max_pooling2d_22 (MaxPoolin (None, 8, 8, 256)      0
g2D)
dropout_27 (Dropout)        (None, 8, 8, 256)      0
conv2d_47 (Conv2D)          (None, 8, 8, 512)    1188160
conv2d_48 (Conv2D)          (None, 8, 8, 1024)   4719616
batch_normalization_7 (Batc (None, 8, 8, 1024)    4096
hNormalization)
max_pooling2d_23 (MaxPoolin (None, 4, 4, 1024)      0
g2D)
dropout_28 (Dropout)        (None, 4, 4, 1024)      0
flatten_7 (Flatten)         (None, 16384)          0
dense_42 (Dense)             (None, 32)             524320
dropout_29 (Dropout)        (None, 32)             0
dense_43 (Dense)             (None, 128)            4224
dropout_30 (Dropout)        (None, 128)            0
dense_44 (Dense)             (None, 256)            33024
dropout_31 (Dropout)        (None, 256)            0
dense_45 (Dense)             (None, 512)            131584
dropout_32 (Dropout)        (None, 512)            0
dense_46 (Dense)             (None, 1024)           525312
dense_47 (Dense)             (None, 10)             10250
=====
Total params: 7,521,002
Trainable params: 7,518,954
Non-trainable params: 2,048
```

```
history = model.fit(x_train_scaled, y_train, epochs=200,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 84.790



Se modifica la arquitectura y se añaden cuatro capas de dropout de 0.3 tras el flatten, una entre cada capa dense, la idea es que se reduzca el overfitting, se amplían las epochs hasta 200 para contrarrestar una posible pérdida de precisión al pasar las imágenes por menos neuronas. La accuracy se ha mantenido muy similar y el overfitting ha disminuido, tal y como se puede apreciar en los gráficos

# VERSIÓN 13

```
model = ks.Sequential()

model.add(ks.layers.Conv2D(64, (3, 3), strides=1, activation='relu',
                           padding='same', input_shape=(32,32,3)))
model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(128, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(256, (3,3), padding='same', activation='relu'))
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Conv2D(512, (3,3), padding='same', activation='relu'))
model.add(ks.layers.Conv2D(1024, (3,3), padding='same', activation='relu'))
model.add(ks.layers.BatchNormalization())
model.add(ks.layers.MaxPooling2D((2, 2)))
model.add(ks.layers.Dropout(0.4))

model.add(ks.layers.Flatten())
model.add(ks.layers.Dense(32, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(128, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(256, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(512, activation='relu'))
model.add(ks.layers.Dropout(0.3))
model.add(ks.layers.Dense(1024, activation='relu'))

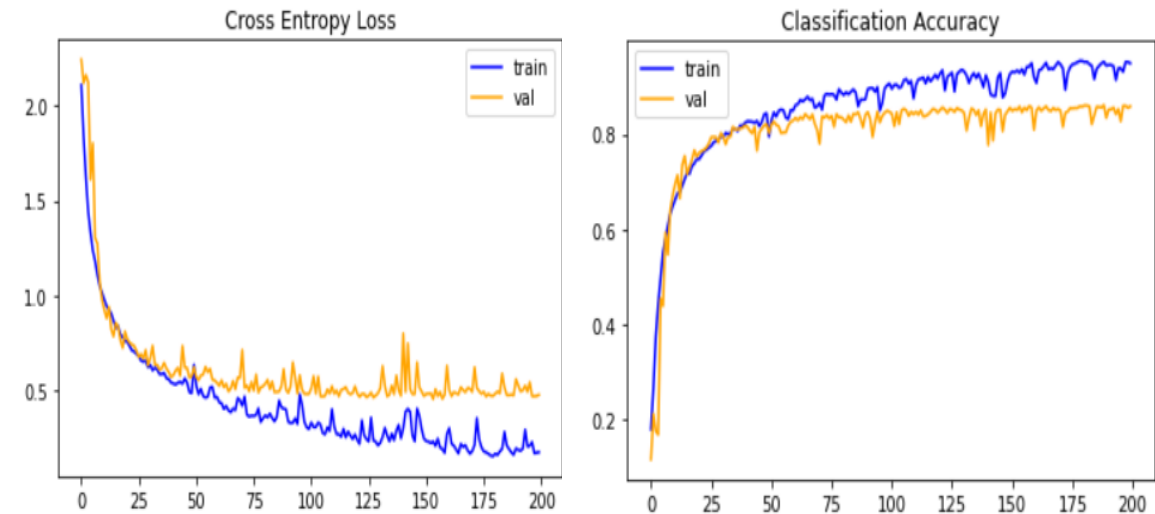
model.add(ks.layers.Dense(10, activation='softmax'))
```

```
model.summary()
=====
conv2d_6 (Conv2D)          (None, 32, 32, 64)        1792
conv2d_7 (Conv2D)          (None, 32, 32, 128)       73856
max_pooling2d_3 (MaxPooling (None, 16, 16, 128)        0
2D)
dropout_7 (Dropout)        (None, 16, 16, 128)       0
conv2d_8 (Conv2D)          (None, 16, 16, 128)       147584
conv2d_9 (Conv2D)          (None, 16, 16, 256)       295168
max_pooling2d_4 (MaxPooling (None, 8, 8, 256)          0
2D)
dropout_8 (Dropout)        (None, 8, 8, 256)        0
conv2d_10 (Conv2D)         (None, 8, 8, 512)         1180160
conv2d_11 (Conv2D)         (None, 8, 8, 1024)        4719616
batch_normalization_1 (Batc (None, 8, 8, 1024)        4096
hNormalization)
max_pooling2d_5 (MaxPooling (None, 4, 4, 1024)          0
2D)
dropout_9 (Dropout)        (None, 4, 4, 1024)       0
flatten_1 (Flatten)        (None, 16384)             0
dense_6 (Dense)            (None, 32)                524320
dropout_10 (Dropout)       (None, 32)                0
dense_7 (Dense)            (None, 128)               4224
dropout_11 (Dropout)       (None, 128)               0
dense_8 (Dense)            (None, 256)               33024
dropout_12 (Dropout)       (None, 256)               0
dense_9 (Dense)            (None, 512)               131584
dropout_13 (Dropout)       (None, 512)               0
dense_10 (Dense)           (None, 1024)              525312
dense_11 (Dense)           (None, 10)                10250
=====
Total params: 7,650,986
Trainable params: 7,648,938
Non-trainable params: 2,048
```

```
history = model.fit(x_train_scaled, y_train, epochs=200,
                    callbacks=[model_checkpoint_callback, earlystopping_loss_callback, earlystopping_acc_callback],
                    batch_size= 512, validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 85.480



Se aumenta el número de neuronas de las dos primeras capas de convoluciones con la idea de aumentar la accuracy del modelo, el resto se mantiene igual. La accuracy ha mejorado hasta situarse en 85.480 pero los gráficos presentan muchos picos.